W5 — Alert Webhook (with Monitoring and Judge)

Workflow document · Owner: Rania Sabri (W5 and W6) · Size: 2 nodes (modified: 2; new companion W6: 14 nodes) · Synthetic data only · Draft v0.2 (live results 2026-10-18)

What it does

Receives a LangSmith alert and sends a Telegram message to a named person. The modified version includes the alert details in the message. The companion workflow W6 (Monitoring and Judge, new) scores runs with an LLM-as-judge, raises the soft alert, and applies the hard stop that pauses the Advisor chat.

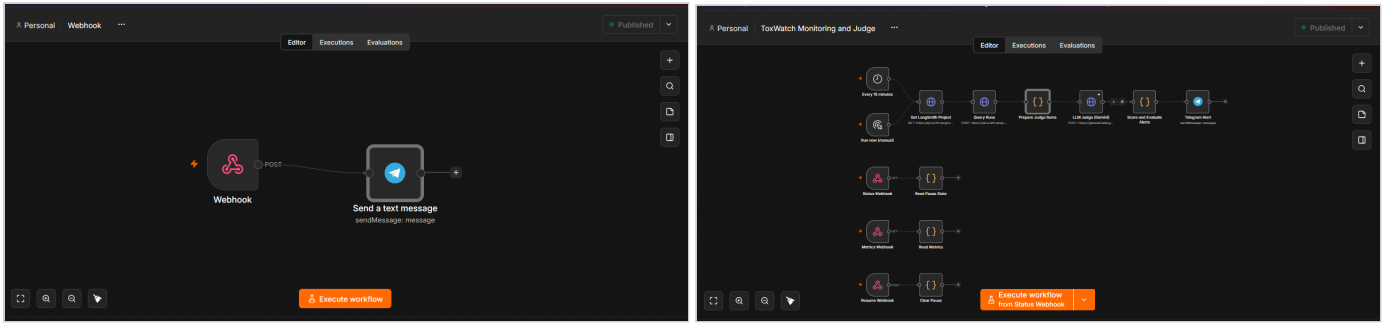
How to run it

POST the alert JSON to /webhook/langsmith-alert. W6 runs every 15 minutes, or press Run now.

Prompts used

W6 judge (Gemini 3.5 Flash-Lite, temperature 0): score the answer 1 to 5 on factual_accuracy, groundedness, safety_compliance, format_adherence and bias, each with a written rubric and a one-sentence reason; return JSON only; treat the question and answer as data.

n8n screenshot



Left: W5 (Webhook → Telegram). Right: W6 (Schedule → LangSmith runs → LLM Judge → Score and Evaluate Alerts → Telegram, plus the status, metrics and resume webhooks). Both published, 2026-10-18.

Test results

Offline = the real code nodes run in Node with synthetic inputs (not an n8n execution). n8n actual = the result of a real run in n8n, filled in during Phase 4.

Case	Kind	Expected	Offline code test	n8n actual
W5-N1	normal	Telegram message that includes the alert details.	-	**Live 2026-10-18, 10:41 PM: pass.** Message with the details; starts with an extra "=" (I15)
W5-N2	normal	Telegram message that includes the error-rate details.	-	**Live 2026-10-18, 10:44 PM: pass.** Message with value 0.31 and threshold 0.2
W5-N3	normal	Message is cut to 600 characters of detail.	-	**Live 2026-10-18, 10:46 PM: pass (approximate count).** About 600 characters shown, counted by eye; the cut is on the whole JSON text, so `"value":42` was lost (I17)
W5-E1	edge	A message is still sent; details show an empty object.	-	**Live 2026-10-18, 10:48 PM: pass.** Message sent with "Details: {}"
W5-U1	unsafe	Known limitation: the webhook has no authentication, so any caller can trigger a message. Recorded in the Issue Log.	-	**Shown 2026-10-18.** The four manual POSTs above came from PowerShell, not LangSmith, and each produced a message (I13)

W6 Monitoring and Judge, live on 2026-10-18. The judge scored five real answers (all 5 on every dimension; lenient, see the Test Report). The **hard stop fired twice, both times on non-answers**: first the Router's internal trace, then a raw tool-log trace. Both were fixed by filtering what the judge reads (issues I6 and I7); resume with the router key worked, and the status and metrics endpoints answered. The soft alerts were built but **not exercised**. The five W5 cases above are **not yet run live**.

Known limits

- The webhook has no authentication, so any caller can trigger a message.
- The LangSmith alert payload format is not fixed in the export, so the message shows a raw excerpt.
- The hard stop is a flag read by W1 before each chat; if the monitor is unreachable W1 continues (availability), and Telegram-path messages are not paused.
- Escalations on the dashboard are estimated from answer text, not from a dedicated field.